

Recommendation System - Amazon Book Review Analysis

Business Understanding



- The days of customers walking into a shop to buy what they need/want are long behind us and worse still if these are items are not basic needs. More and more clients prefer to make purchases from the comfort of their home.
- The goods that a retailer is able to market online is limitless however customers easily get tired of scrolling through an endless catalogue of items for sale.
- Therefore rises the need for a recommendation system that will enable a client have a seamless buying experience. The reading culture is changing hence our choice of the amazon books dataset.
- A recommendation system will enable buyers get the most ideal and trending books to buy.
- The target audience would be both the retailers and the purchasers.

Data Understanding & Source

- The data has been obtained from <https://amazon-reviews-2023.github.io/> (<https://amazon-reviews-2023.github.io/>) and in jsonl format. An efficient format for storing data that is unstructured or produced over time.
- It contains a list of books sold in Amazon. The original dataset contains 4 million rows, from 1996 to 2023. We trimmed it to 300k rows from the year 2023 to make it easier to work with.

- Vital information was missing from the dataset (price, book title). This was obtained by merging the data with the metadata dataset.
- The data contains following features/columns in the dataset.

Column Name	Description
rating	Rating of the product (from 1.0 to 5.0).
title_x	Title of the user review.
text	Text body of the user review.
images	Links to images (comma-separated if multiple).
asin(product key)	Unique identifier for the product.
parent_asin	Identifier for the parent product (applicable for variations).
user_id	Unique identifier for the reviewer.
timestamp	Date and time of the review.
helpful_vote	Number of helpful votes received by the review.
verified_purchase	Indicates whether the reviewer purchased the product (True/False).
main_Category	Main category (domain) to which the product belongs (e.g., Electronics, Clothing).
title_y	Name of the product as mentioned in the review.

Data Importation

```
In [ ]: 1 # Mount the google drive
        2 # from google.colab import drive
        3 # drive.mount('/content/drive')
```

```
In [ ]: 1 # import the necessary libraries
        2
        3 import pandas as pd
        4 import json
        5 import numpy as np
        6 import matplotlib.pyplot as plt
        7 import seaborn as sns
```

```

In [ ]: 1 # Load the merged dataset
2 file_path = '/content/drive/MyDrive/merged_Books.jsonl'
3 # Initialize an empty list to store the parsed JSON objects
4 data = []
5
6 # Read each Line of the JSON Lines file and parse it
7 with open(file_path, 'r') as f:
8     for line in f:
9         data.append(json.loads(line))
10
11 # Convert the List of JSON objects into a DataFrame
12 df = pd.DataFrame(data)
13 df.head()

```

```

Out[5]:

```

	rating	title_x	text	images	asin	parent_asin
0	5	Wonderful and Inspiring	This book is wonderful and inspiring for kids ...	[]	B0C6Z8N9N8	B0C6Z8N9N8 AG2FEEHWHCQELOHBIDQC
1	5	Awesome book	This is a wonderful children's book! My daught...	[]	B0C6Z8N9N8	B0C6Z8N9N8 AERUMG7KTKZAIQ3PO5I
2	5	Amazing	Product arrived quickly in great condition. Be...	[]	1401241883	1401241883 AEK3AFSE3D2BSOC6XI65
3	5	Got this at a great price.	I payed \$89.00 dollars. When it first came out...	[]	1401241883	1401241883 AFPYBFVIJI3GFPPFANF
4	5	The Best of the Best	Neil Gaimans stories are spellbinding. Moreove...	[]	1401241883	1401241883 AECBBBUARXJEZYZS2PXI

Data Understanding

```
In [ ]: 1 # Preview the attributes of the data
        2
        3 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300000 entries, 0 to 299999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rating                300000 non-null  int64
1   title_x               300000 non-null  object
2   text                  300000 non-null  object
3   images                300000 non-null  object
4   asin                  300000 non-null  object
5   parent_asin           300000 non-null  object
6   user_id               300000 non-null  object
7   timestamp             300000 non-null  int64
8   helpful_vote          300000 non-null  int64
9   verified_purchase     300000 non-null  bool
10  main_category          299989 non-null  object
11  title_y               300000 non-null  object
12  price                  267594 non-null  object
dtypes: bool(1), int64(3), object(9)
memory usage: 27.8+ MB
```

```
In [ ]: 1 # Review the rows and columns of the data
        2 df.shape
```

Out[7]: (300000, 13)

```
In [ ]: 1 # Min and max rating
        2 print('Max rating:', df['rating'].max())
        3 print('Min rating:', df['rating'].min())
```

```
Max rating: 5
Min rating: 1
```

- Some book titles appear multiple times. We can get the value counts for the most frequent.

```
In [ ]: 1 # Most frequent boooks
        2 df['title_y'].value_counts()
```

```
Out[9]: title_y
The Sacrifice: A Dark Revenge Romance
476
Spare
463
The Maid's Diary: A Novel
354
The Serpent and the Wings of Night (Crowns of Nyaxia Book 1)
296
Stone Maidens
260

...
The Long Goodbye: A Philip Marlowe Novel, Book 6
1
Corpse in the Mead Hall: A Viking Witch Cozy Mystery (The Viking Witch Cozy M
ysteries Book 6)      1
Metro Wine Map of France
1
Souffle Cookbook: Souffle Recipes from Around the World: Souffle Cookbook For
You      1
Brut Y Brenhinedd
1
Name: count, Length: 156087, dtype: int64
```

- Getting the value counts for users that give multiple review ('user_id')

```
In [ ]: 1 # Top 5 users based on the number of ratings
        2 top_five_users = df.groupby('user_id').size().sort_values(ascending=False)
        3 top_five_users
```

```
Out[10]: user_id
AHK67LFXJBYE5APXUTYTJTDSHL4A      264
AGWMG5ARMSS5U2FMSSMPNML6MTNQ_1    121
AGVBYI2T5QRJVZ6KX2YH7LHF7YRQ      90
AFMBF3NCA6H2AH06D2WX6SUBPELA      86
AENPLYFNCNXWGB3XF2HPD5EKJD6Q      80
dtype: int64
```

```
In [ ]: 1 # Visual for book title appearing more than 200 times
2 filtered_counts = df['title_y'].value_counts()[lambda x: x >= 100]
3 # Increase plot width using figsize
4 plt.figure(figsize=(8, 10))
5
6 # Plot the distribution
7 filtered_counts.plot(kind='barh')
8 plt.show()
```



```
In [ ]: 1 # Books only appearing once and more than 50 times in the data
2 single_mention = df['title_y'].value_counts()[lambda x: x == 1]
3 above_fifty_mention = df['title_y'].value_counts()[lambda x: x >= 50]
4 print(len(above_fifty_mention))
5 print(len(single_mention))
```

180

114626

```
In [ ]: 1 # Having a df with unique book titles
        2 df = df.drop_duplicates(subset=['title_y'])
        3 df['title_y'].value_counts()
```

```
Out[13]: title_y
Irish Rain
1
Of Life: The Rollercoaster
1
The Sandman Omnibus Vol. 1
1
The Keeper of Happy Endings
1
The Echo of Old Books: A Novel
1
..
123 Counting Sticker Book (My Little World)
1
Soap Making Business Startup: How to Start, Run & Grow a Million Dollar Success From Home!
1
How to Write Dazzling Dialogue: The Fastest Way to Improve Any Manuscript (Bell on Writing)
1
An American Demon: A Memoir
1
Ori: The Ultimate Guide to Spiritual Intuition, Yoruba, Odu, Egbe, Orishas, and Ancestral Veneration (African Spirituality)
1
Name: count, Length: 156087, dtype: int64
```

```
In [ ]: 1 pip install wordcloud
```

Collecting wordcloud

Downloading wordcloud-1.9.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (511 kB)

511.1/511.1 kB 2.2 MB/s eta 0:00

Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.10/dist-packages (from wordcloud) (1.25.2)

Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from wordcloud) (10.3.0)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from wordcloud) (3.7.1)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.2.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (4.51.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.4.5)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (24.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.16.0)

Installing collected packages: wordcloud

Successfully installed wordcloud-1.9.3

In []:

```
1 # Create a wordcloud of the most purchased books
2 from wordcloud import WordCloud
3
4 # Count occurrences of each book title
5 book_counts = df['title_y'].value_counts()
6
7 # Select the top 50 most frequent books
8 top_50_books = book_counts.head(50)
9
10 # Convert to a dictionary where the keys are the book titles and the value
11 word_freq = top_50_books.to_dict()
12
13 # Generate the word cloud
14 wordcloud = WordCloud(width=1000, height=600, background_color='white').ge
15
16 # Display the word cloud
17 plt.figure(figsize=(10, 10))
18 plt.imshow(wordcloud, interpolation='bilinear')
19 plt.axis('off')
20 plt.show()
```



- Visualizing distribution of ratings.

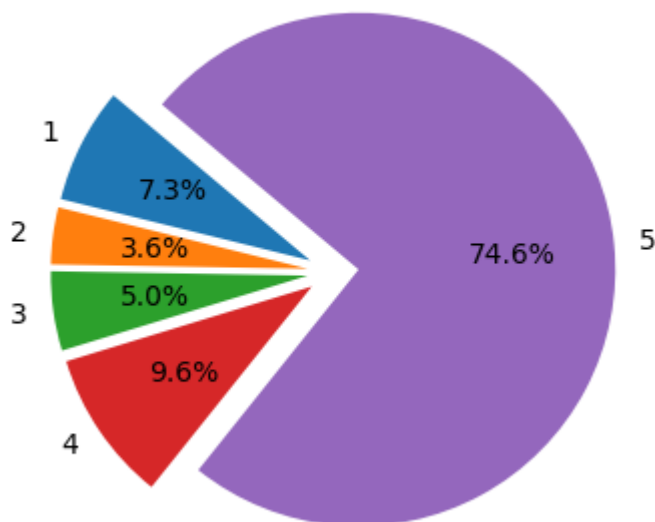
```

In [ ]: 1 # Group sales data rating by count of title
        2
        3 rating_df = df[['rating', 'title_y']].copy().groupby('rating').count()
        4 print(rating_df)
        5
        6 # Explode settings
        7 explode = (0.1, 0.1, 0.1, 0.1, 0.1) # Explode all the slices
        8
        9 # Plot
       10 plt.figure(figsize=(4, 4))
       11 plt.pie(rating_df['title_y'], labels=rating_df.index, autopct='%1.1f%%', s
       12 plt.axis('equal')
       13 plt.title('Distribution of Ratings')
       14 plt.show()

```

	title_y
rating	
1	11341
2	5578
3	7781
4	14911
5	116476

Distribution of Ratings



- Books with a rating of 5 take up 74.6% of the data. This necessitates for more features to be used in recommendation, that is title for rating and rating text.

```
In [ ]: 1 # Books rating count
        2 rating_count = df['rating'].value_counts()
        3 rating_count
```

```
Out[17]: rating
5      116476
4       14911
1       11341
3        7781
2         5578
Name: count, dtype: int64
```

```
In [ ]: 1 # Get user count information
        2 user_count = df['user_id'].value_counts()
        3 single_user_mention = df['user_id'].value_counts()[lambda x: x == 1]
        4 print('Number of users appearing once:', len(single_user_mention) )
        5 print('Number of unique users:', len(user_count))
```

```
Number of users appearing once: 105282
Number of unique users: 122457
```

Exploring the 'main_category' field

```
In [ ]: 1 # Identify the unique values in the
        2
        3 print(df['main_category'].unique())
```

```
['Books' 'Buy a Kindle' 'Musical Instruments' 'Audible Audiobooks' ''
 'Toys & Games' 'Office Products' 'AMAZON FASHION' 'Amazon Home' None
 'Tools & Home Improvement' 'Arts, Crafts & Sewing'
 'Industrial & Scientific']
```

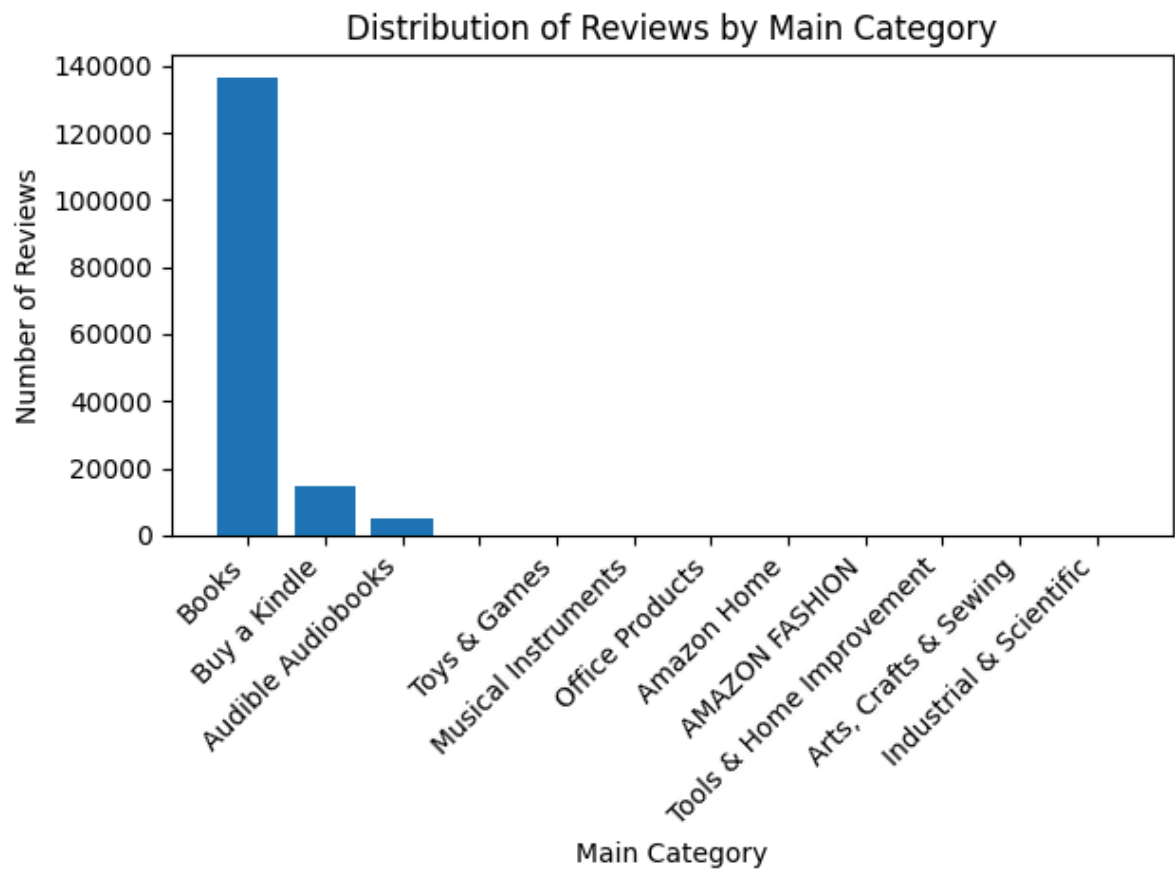
```
In [ ]: 1 # Identify the unique values in the 'main_category'
        2 print(df['main_category'].value_counts())
```

```
main_category
Books      136261
Buy a Kindle    14781
Audible Audiobooks    4872
              132
Toys & Games    15
Musical Instruments    5
Office Products    5
Amazon Home    5
AMAZON FASHION    1
Tools & Home Improvement    1
Arts, Crafts & Sewing    1
Industrial & Scientific    1
Name: count, dtype: int64
```

```

In [ ]: 1 # Count the occurrences of each category
2 category_counts = df['main_category'].value_counts()
3
4 # Create a bar chart
5 plt.bar(category_counts.index, category_counts.values)
6 plt.xlabel('Main Category')
7 plt.ylabel('Number of Reviews')
8 plt.title('Distribution of Reviews by Main Category')
9 plt.xticks(rotation=45, ha='right')
10
11 # Show the plot
12 plt.tight_layout()
13 plt.show()

```



- From the plot above we see that most of the books are classified in the **Books**, **Buy a Kindle** and **Audiobooks** category. Other categories do not have more than 30 books, and 182 of them not categorized.
- Category column does not give the book genres accurately and therefor will be dropped

Data Cleaning

- First we drop the columns not needed
 - images
 - asin
 - parent_asin

- timestamp
- verified_purchase

```
In [ ]: 1 # Drop the columns
2 columns_to_drop = ['images', 'asin', 'parent_asin', 'timestamp', 'verified_
3 df1 = df.drop(columns_to_drop, axis=1)
4
5 df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 156087 entries, 0 to 299999
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   rating          156087 non-null  int64
1   title_x         156087 non-null  object
2   text            156087 non-null  object
3   user_id         156087 non-null  object
4   helpful_vote    156087 non-null  int64
5   main_category   156080 non-null  object
6   title_y         156087 non-null  object
7   price           139540 non-null  object
dtypes: int64(2), object(6)
memory usage: 10.7+ MB
```

```
In [ ]: 1 # Remove the words 'from' and 'None' from the price column
2 df1['price'] = df1['price'].astype(str).str.replace(r'(from|None)\s*', '',
3
4 # Remove special characters from the price column
5 df1['price'] = df1['price'].replace(['', '-'], np.nan)
6
7 # Convert the price column to data type float
8 df1['price'] = df1['price'].astype(float)
9 df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 156087 entries, 0 to 299999
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   rating          156087 non-null  int64
1   title_x         156087 non-null  object
2   text            156087 non-null  object
3   user_id         156087 non-null  object
4   helpful_vote    156087 non-null  int64
5   main_category   156080 non-null  object
6   title_y         156087 non-null  object
7   price           137727 non-null  float64
dtypes: float64(1), int64(2), object(5)
memory usage: 10.7+ MB
```

```
In [ ]: 1 # Sum up the null values in the price column
        2
        3 df1['price'].isnull().sum()
```

Out[24]: 18360

```
In [ ]: 1 # Fill null values in 'price' column with the mean
        2
        3 df1['price'] = df1['price'].fillna(df1['price'].mean())
        4 df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 156087 entries, 0 to 299999
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   rating           156087 non-null  int64
1   title_x          156087 non-null  object
2   text             156087 non-null  object
3   user_id          156087 non-null  object
4   helpful_vote     156087 non-null  int64
5   main_category    156080 non-null  object
6   title_y          156087 non-null  object
7   price            156087 non-null  float64
dtypes: float64(1), int64(2), object(5)
memory usage: 10.7+ MB
```

- With no missing values, we move to renaming the columns to more meaningful titles

```
In [ ]: 1 # Rename the title_x and title_y column to title_rating and title_book res
        2
        3 df1 = df1.rename(columns={'title_x': 'title_rating', 'title_y': 'title_book'})
        4 df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 156087 entries, 0 to 299999
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   rating           156087 non-null  int64
1   title_rating     156087 non-null  object
2   text             156087 non-null  object
3   user_id          156087 non-null  object
4   helpful_vote     156087 non-null  int64
5   main_category    156080 non-null  object
6   title_book       156087 non-null  object
7   price            156087 non-null  float64
dtypes: float64(1), int64(2), object(5)
memory usage: 10.7+ MB
```

```
In [ ]: 1 # Preview the data with new columns
        2 df1.head(3)
```

Out[27]:

	rating	title_rating	text	user_id	helpful_vote	main_categc
0	5	Wonderful and Inspiring	This book is wonderful and inspiring for kids ...	AG2FEEHWHCQELOHBIDQDROZ3LSNA	0	Boc
2	5	Amazing	Product arrived quickly in great condition. Be...	AEK3AFSE3D2BSOC6XI65XNO23MKQ	0	Boc
5	5	Good	Just what I expectef	AFAHBYOMYBR5JNAYCR5P2PCUBMWQ	0	Boc

```
In [ ]: 1 # Convert text and title_rating column to lower case and remove punctuatio
        2 import string
        3
        4 def clean_text(text):
        5     if isinstance(text, str):
        6         text = text.lower()
        7         # Remove punctuation marks
        8         translator = str.maketrans('', '', string.punctuation)
        9         return text.translate(translator)
        10    else:
        11        return str(text)
        12
        13 df1['text'] = df1['text'].apply(lambda x: clean_text(x))
        14 df1['title_rating'] = df1['title_rating'].apply(lambda x: clean_text(x))
```

In []:

```
1 # Tokenize and remove stop words from the text and title_rating columns
2 import nltk
3 from nltk.corpus import stopwords
4 nltk.download('punkt')
5 nltk.download('stopwords')
6 from nltk.tokenize import word_tokenize
7
8 stop_words = set(stopwords.words('english'))
9
10
11 def remove_stopwords(text):
12     """
13     This function removes stop words from a given text string.
14
15     Args:
16         text (str): The string to remove stop words from.
17
18     Returns:
19         list: A list of words after removing stop words from the original text
20     """
21     # Tokenize input text
22     tokens = word_tokenize(text)
23     # filter out stop words and return list of words without stop words
24     filtered_tokens = [word for word in tokens if word not in stop_words]
25     return filtered_tokens
26
27 # Apply the remove_stopwords function to the 'text' and 'title_rating'
28 df1['tokenized_text'] = df1['text'].apply(lambda x: remove_stopwords(x))
29 df1['tokenized_title_rating'] = df1['title_rating'].apply(lambda x: remove_stopwords(x))
30 df1.sample(3)
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

Out[29]:

	rating	title_rating	text	user_id	helpful_vote	main_
172507	5	informative	enjoyed reading it is informative with nice pi...	AHITNFJGUWWYDTF46K6RAAE7I6GQ	1	
104228	5	fantastic read	absolutely loved this book fantastic story and...	AF7EZ3WWX2FBB7ATI2S6BKGPGGCA	0	
272084	4	essentially a shorter version of other thinkin...	i got this and other grade 3 curriculum books ...	AERUHOWCRKOJYLZZ2RARPFFA2CXA	0	


```

In [ ]: 1 # Display a frequency distribution of the most common words
        2
        3 from nltk.probability import FreqDist
        4 from itertools import chain
        5
        6 def common_words(df, column, n=15):
        7     all_tokens = list(chain.from_iterable(df1['tokenized_title_rating']))
        8     fdist = FreqDist(all_tokens)
        9     return fdist.most_common(n)
       10
       11 common_words(df1, 'tokenized_title_rating', 15)

```

```

Out[30]: [('book', 27586),
          ('great', 19943),
          ('read', 10661),
          ('good', 9299),
          ('love', 6119),
          ('story', 5379),
          ('excellent', 3931),
          ('fun', 3845),
          ('', 3432),
          ('beautiful', 2918),
          ('amazing', 2875),
          ('cute', 2577),
          ('series', 2376),
          ('best', 2371),
          ('perfect', 2297)]

```

```

In [ ]: 1 # Add 'book' and ' to the stop words list and remove them from the tokeniz
        2
        3 additional_stop_words = {'book', '', 'story'}
        4
        5 stop_words.update(additional_stop_words)
        6
        7 df1['tokenized_title_rating'] = df1['title_rating'].apply(lambda x: remove
        8
        9 print(common_words(df1, 'tokenized_title_rating', 10))
       10

```

```

[('great', 19943), ('read', 10661), ('good', 9299), ('love', 6119), ('excellent', 3931), ('fun', 3845), ('beautiful', 2918), ('amazing', 2875), ('cute', 2577), ('series', 2376)]

```

```
In [ ]: 1 # Filter the DataFrame to include only rows where the title rating contain
2 placeholder_word = 'great'
3
4 common_words_books_df = df1[df1['tokenized_title_rating'].apply(lambda x:
5 common_words_books_count_df = common_words_books_df['title_book'].value_co
6 print(common_words_books_count_df)
7
8 # Select the top 50 books with the most common word
9 top_50_common_words_books_df = common_words_books_count_df.head(50)
10
11 # Convert into a dictionary where the keys are the book titles and the val
12 word_freq_great = top_50_common_words_books_df.to_dict()
13
14 # Generate a word cloud
15 wordcloud = WordCloud(width=800, height=400, background_color='white').gen
16
17 # Display the word cloud using matplotlib
18 plt.figure(figsize=(10, 5))
19 plt.imshow(wordcloud, interpolation='bilinear')
20 plt.axis('off')
21 plt.show()
```

title_book

Fido's Magical Quest: An Adventure for All Ages

1

THE RIVER'S EDGE a gripping crime thriller full of stunning twists (JACKMAN & EVANS Book 10)

1

Dead Fall: A Thriller (The Scot Harvath Series Book 22)

1

The Megalodon Mix-Up (A Charlie Rhodes Cozy Mystery Book 4)

1

123 Counting Sticker Book (My Little World)

1

..

Disney Before the Story: Elsa's Icy Rescue

1

A Fire Sparkling

1

The "I Love My Air Fryer" 5-Ingredient Recipe Book: From French Toast Sticks to Buttermilk-Fried Chicken Thighs, 175 Quick and Easy Recipes ("I Love My" Cookbook Series) 1

Losing Hope: A Novel (2) (Hopeless)

1

The Milkmaid: The Royal Betrayal: Book One

1

Name: count, Length: 19621, dtype: int64

Weight Loss Journal for Women: Simple 10-Week Food and Fitness Logbook / Motivational Diet and Exercise Planner with Daily Health and Beauty Affirmations / Weight Loss and Fitness Tracker

The Megalodon Mix-Up (A Charlie Rhodes Cozy Mystery Book 4)

Toddler's First Coloring Book Ages 1-4: Amazing Coloring Sheets for Toddlers and Kids Ages 1, 2, 3 & 4 | Over 100 Simple pictures for coloring, doodling ... Fruits and everyday objects to color for kids
Captain's Log (Scifi)

THE RIVER'S EDGE a gripping crime thriller full of stunning twists (JACKMAN & EVANS Book 10)

Perfect Sous Vide At Home Cookbook: Compatible with Anova & Most Sous Vide Cookers - 101 Restaurant-Quality Recipes Your Family Will Love!
Rising Rain Wild Card (A Sinister Thriller) Just a Kid from Swampoodle to Vietnam

Emeril Lagasse Power Air Fryer 360 Cookbook 2022: The Ultimate Everyday Delicious Days of Power Air Fryer 360 Recipes Dead Fall: A Thriller (The Scot Harvath Series Book 22)

The Heirloom Garden: Traditional Plants and Skills for the Modern World

Huber's Journal Playing the Odds: The MacGregors, Book 1

Can't Hurt Me: Master Your Mind and Defy the Odds

Heir to Storm and Flame: A totally addictive MM fantasy romance (Court of Broken Bonds)

123 Counting Sticker Book (My Little World)

Writing Gatsby: The Real Story of the Writing of the Greatest American Novel

August Wilson Century Cycle

Somnata: There Once Was a Town Called Tranquility

Frankenstein (Wordsworth Collector's Editions)

Anatomy for the Artist

Ty Beanie's Tracker Third Edition

In My Mind

Captivating Expanded Edition: Unveiling the Mystery of a Woman's Soul

Grendel's Labyrinth (John Decker Supernatural Thrillers Book 4)

Georgia Milestones Assessment System Test Prep: 8th Grade Math Practice Workbook and Full-length Online Assessments: GMAS Study Guide (GMAS by Lunos Learning)

Fido's Magical Quest: An Adventure for All Ages

31 Spooky Writing Prompts for Kids: Growth Mindset Questions | Creative Writing | Opinion Writing | Expository Writing | Narrative Writing
Travels In West Africa

Invisible Monsters: A Novel

The Storyteller: Tales of Life and Music

Adorable Creepy Mushrooms: A cute coloring book with mushroom monsters | Relaxation and fun (for Adults, Teens and Kids) (Cute Little Critters)

1001 All-Natural Secrets to Pest Control (If They Are FLYING CRAWLING BURROWING OR SNEAKING IN THIS BOOK HAS THE SOLUTION)

The Soul of the City: Le Petit Théâtre du Vieux Carré

Four and Ready to Explore: Taking on Chores and Reading Galore!

The Ultimate Mindfulness Coloring Book for Adults: Rediscover Relaxation, Focus, Stress Relief and Your Inner Zen. With amazing patterns, animals and mandalas. 1-sided designs.

The Professional Poker Dealer's Handbook: Expanded Edition

The Psychology of Money: Timeless Lessons on Wealth, Greed, and Happiness

Hear Me Roar: Women, Motorcycles and the Rapture of the Road, New Ed.

Exotic Room's Revelation Summer

Going To Meet the Man

Lake(Plate) Houses: Embracing the Landscape

Self-defense or Jiu-Jitsu achievable by everyone

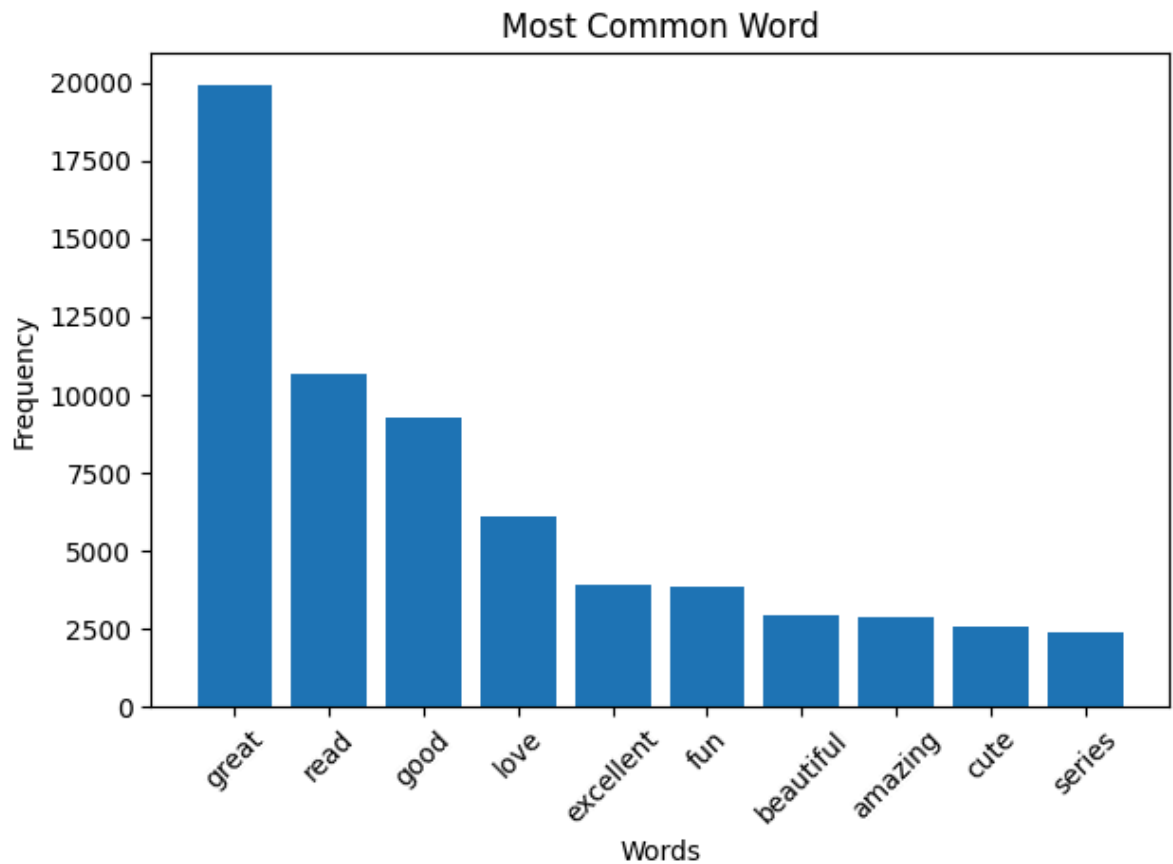
Finance 101 For Kids: The ABC of Money

Emerson's Note (A Large Paperback and New Firstland Mystery Book 28)

My Life Planner

Ancient Civilizations: Lamentations and Magic Book 1

```
In [ ]: 1 # Get the 10 most common words (excluding additional stop words)
2 most_common_words = common_words(df1, 'tokenized_title_rating', 10)
3
4 # Unpack tuples into separate lists for words and counts
5 words, counts = zip(*most_common_words)
6
7 # Create a bar chart to visualize the most common words
8 plt.bar(words, counts)
9 plt.xlabel('Words')
10 plt.ylabel('Frequency')
11 plt.title('Most Common Word')
12 plt.xticks(rotation=45)
13 plt.tight_layout()
14 plt.show()
```



Modelling

```
In [ ]: 1 df1.sample(2)
```

Out[34]:

	rating	title_rating	text	user_id	helpful_vote	main
183793	2	overproduced and underedited	good to have these important poems in book for...	AGNAFS6TZDSJGMRATTLZKGBIPTWQ	0	
285966	5	good read	what a interesting life	AECE5YC3NO3UV67LD4QRNUKSPKMA	0	

Sentiment Analysis:

The goal is to understand the sentiment expressed in the review text (positive or negative). The sentiment analysis here includes traditional machine learning algorithms involving Multinomial Naive Bayes analysis, Support Vector Machines (SVM) and Random Forest.

Multinomial Naive Bayes classifier

TF-IDF

Term Frequency x Inverse Document Frequency

The number of times
a word appears in a document

A measure of whether a term is common or rare in a
collection of documents

```

In [ ]: 1 from sklearn.feature_extraction.text import TfidfVectorizer
2 from sklearn.model_selection import train_test_split
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn.metrics import accuracy_score, classification_report
5
6 # Select the features
7 X = df1['tokenized_title_rating']
8 y = (df1['rating'] > 3).astype(int) # Convert ratings to binary labels (1
9
10 # Convert each review text list to a single string
11 X_str = [' '.join(tokens) for tokens in X]
12
13 # Split data into train and test sets(20%)
14 X_train, X_test, y_train, y_test = train_test_split(X_str, y, test_size=0.
15
16 # Create TF-IDF vectors
17 tfidf_vectorizer = TfidfVectorizer(max_features=5000)
18 X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
19 X_test_tfidf = tfidf_vectorizer.transform(X_test)
20
21 # Train Multinomial Naive Bayes classifier
22 nb_classifier = MultinomialNB()
23 nb_classifier.fit(X_train_tfidf, y_train)
24
25 # Predictions
26 y_pred = nb_classifier.predict(X_test_tfidf)
27
28 # Evaluate performance
29 accuracy = accuracy_score(y_test, y_pred)
30 report = classification_report(y_test, y_pred)
31
32 # Classification report
33 report = classification_report(y_test, y_pred)
34 print(report)

```

	precision	recall	f1-score	support
0	0.83	0.45	0.58	4926
1	0.90	0.98	0.94	26292
accuracy			0.90	31218
macro avg	0.87	0.72	0.76	31218
weighted avg	0.89	0.90	0.89	31218

```
In [ ]: 1 def evaluate_model(feature_count):
2         # Vectorization with the specified number of features
3         tfidf_vectorizer = TfidfVectorizer(max_features=feature_count)
4         X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
5         X_test_tfidf = tfidf_vectorizer.transform(X_test)
6
7         # Train the classifier
8         nb_classifier = MultinomialNB()
9         nb_classifier.fit(X_train_tfidf, y_train)
10
11        # Make predictions
12        y_pred = nb_classifier.predict(X_test_tfidf)
13
14        # Calculate accuracy
15        accuracy = accuracy_score(y_test, y_pred)
16        return accuracy
```

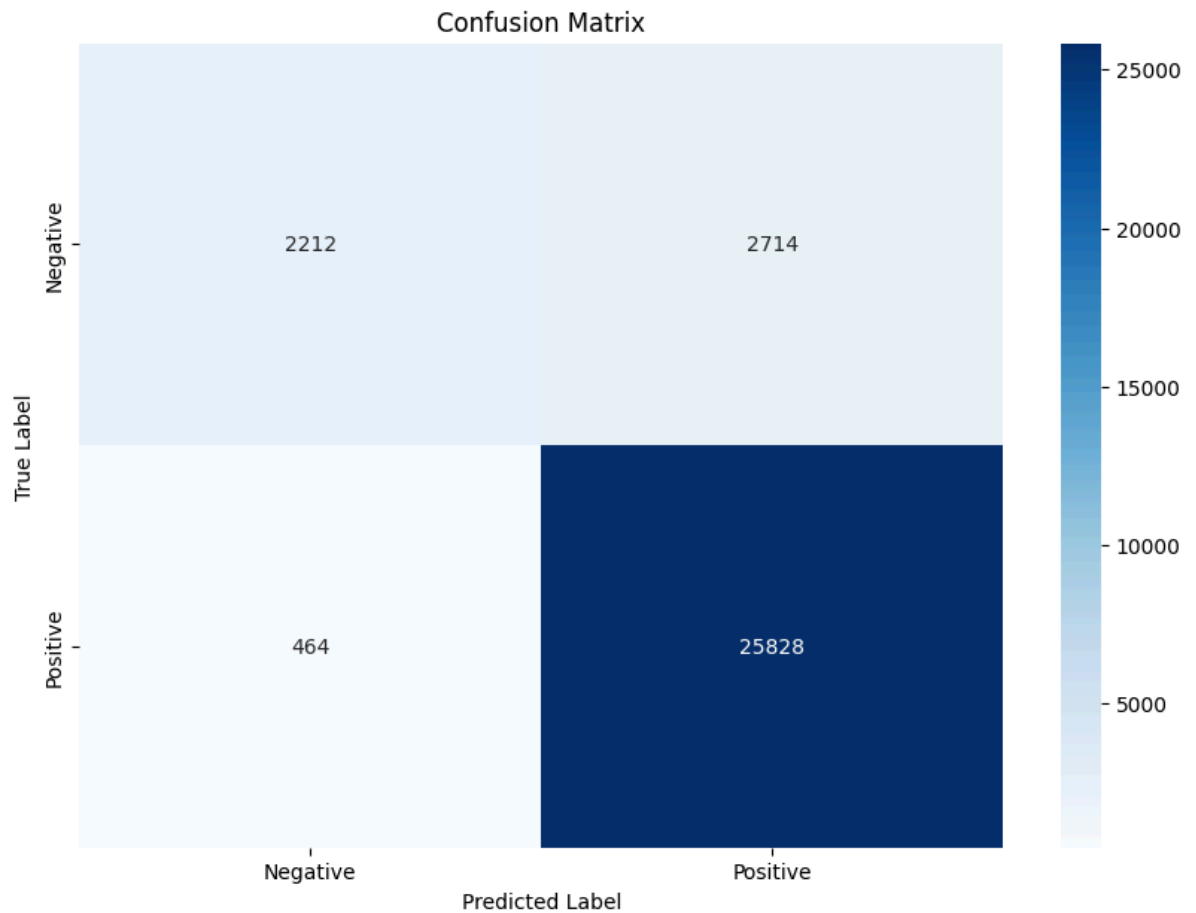
- Create a function to check for best feature count

```
In [ ]: 1 print('Accuracy for 1000 features:',evaluate_model(1000))
2 print('Accuracy for 3000 features:',evaluate_model(3000))
3 print('Accuracy for 5000 features:',evaluate_model(5000))
4 print('Accuracy for 10000:',evaluate_model(10000))
```

```
Accuracy for 1000 features: 0.8838810942405023
Accuracy for 3000 features: 0.8949003779870588
Accuracy for 5000 features: 0.8981997565507079
Accuracy for 10000: 0.8995451342174386
```

The difference between 5000 features and 10000 features is 0.001 which we consider as negligible and proceed with using 5000.

```
In [ ]: 1 # Generate confusion matrix
2 from sklearn.metrics import confusion_matrix
3
4 conf_matrix = confusion_matrix(y_test, y_pred)
5
6 # Plot confusion matrix using seaborn
7 plt.figure(figsize=(10, 7))
8 sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['
9 plt.title('Confusion Matrix')
10 plt.xlabel('Predicted Label')
11 plt.ylabel('True Label')
12 plt.show()
```



- Overall the model performed well with an accuracy of 90%. This translates to 28040 accurate reviews and compared to 3178 negative predictions. This led to selecting a different model; Support Vector Machines that can handle imbalanced data better.

SVM

```
In [ ]: 1 from sklearn.svm import LinearSVC
        2
        3 # Create TF-IDF vectors with 5000 max features
        4 tfidf_vectorizer = TfidfVectorizer(max_features=5000)
        5 X_tfidf = tfidf_vectorizer.fit_transform(X_str)
        6
        7 # Split data into train and test sets
        8 X_train, X_test, y_train, y_test = train_test_split(X_tfidf, y, test_size=
        9
        10 # Train LinearSVC classifier
        11 svm_classifier = LinearSVC(dual=False)
        12 svm_classifier.fit(X_train, y_train)
        13
        14 # Predictions
        15 y_pred = svm_classifier.predict(X_test)
        16
        17 # Evaluate performance
        18 accuracy = accuracy_score(y_test, y_pred)
        19 report = classification_report(y_test, y_pred)
        20 print(report)
```

	precision	recall	f1-score	support
0	0.79	0.50	0.62	4926
1	0.91	0.98	0.94	26292
accuracy			0.90	31218
macro avg	0.85	0.74	0.78	31218
weighted avg	0.89	0.90	0.89	31218

Random Forest

```
In [ ]: 1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.model_selection import cross_val_predict
3 from sklearn.metrics import classification_report
4
5 # Train Random Forest classifier
6 rf_classifier = RandomForestClassifier(n_estimators=50, n_jobs=-1, max_dep
7
8 # Fit Random Forest classifier
9 rf_classifier.fit(X_train_tfidf, y_train)
10
11 # Perform cross-validation and make predictions
12 y_pred_cv = cross_val_predict(rf_classifier, X_tfidf, y, cv=5)
13
14 # Evaluate performance using classification report
15 report = classification_report(y, y_pred_cv)
16 print(report)
```

	precision	recall	f1-score	support
0	0.95	0.09	0.16	24700
1	0.85	1.00	0.92	131387
accuracy			0.86	156087
macro avg	0.90	0.54	0.54	156087
weighted avg	0.87	0.86	0.80	156087

- First we assess the model's ability to generalize. We perform a grid search with cross-validation to find the optimal max_depth value

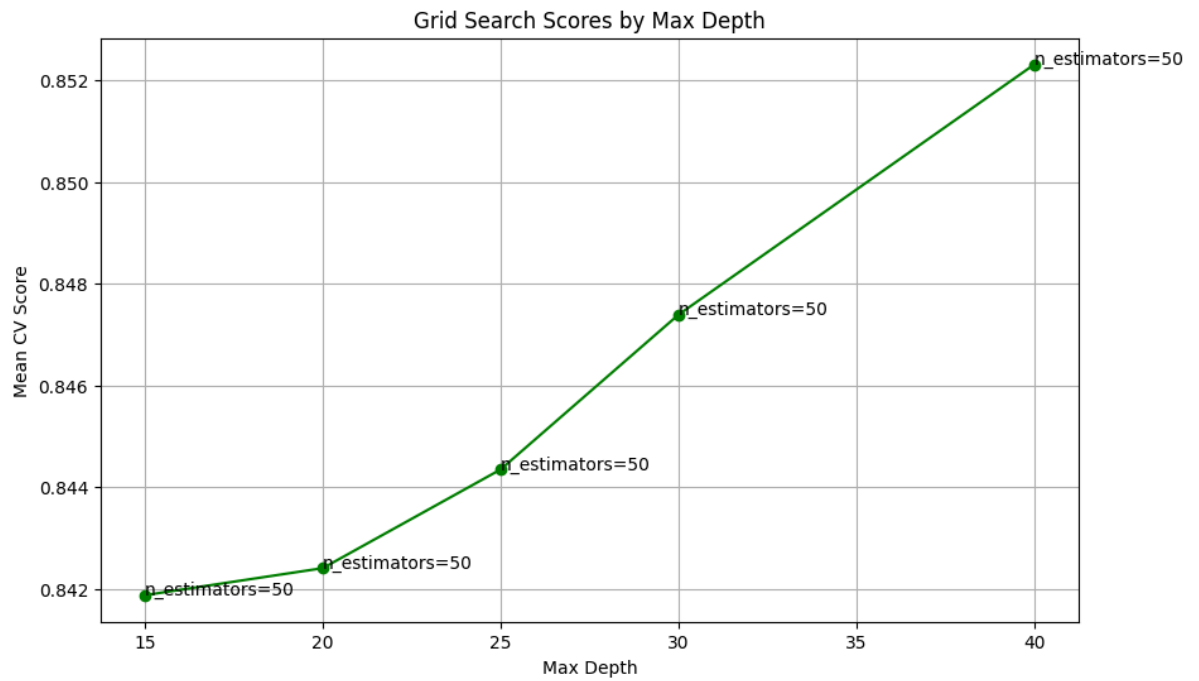
```

In [ ]: 1 from sklearn.model_selection import GridSearchCV
2
3 # Define the parameter grid
4 param_grid = {
5     'max_depth': [15, 20, 25, 30, 40]
6 }
7
8 # Initialize the classifier
9 rf_classifier = RandomForestClassifier(n_estimators=50, n_jobs=-1)
10
11 # Initialize GridSearchCV
12 grid_search = GridSearchCV(rf_classifier, param_grid, cv=5, scoring='accuracy')
13
14 # Perform grid search
15 grid_search.fit(X_tfidf, y)
16
17 # Get the best parameters and the corresponding score
18 best_params = grid_search.best_params_
19 best_score = grid_search.best_score_
20
21 print(f"Best parameters: {best_params}")
22 print(f"Best cross-validation score: {best_score}")
23
24 # Visualize
25
26 # Extract the mean test scores for each parameter setting in the grid
27 mean_scores = grid_search.cv_results_['mean_test_score']
28
29 # Number of trees in random forest for each grid search iteration
30 num_trees = [50] * len(param_grid['max_depth'])
31
32 # Max depth values for each grid search iteration
33 max_depth_values = param_grid['max_depth']
34
35 # Plotting the scores
36 plt.figure(figsize=(10, 6))
37 plt.plot(max_depth_values, mean_scores, marker='o', linestyle='--', color='blue')
38
39 # Annotating the number of trees for each point
40 for i, txt in enumerate(num_trees):
41     plt.annotate(f'n_estimators={txt}', (max_depth_values[i], mean_scores[i]))
42
43 plt.xlabel('Max Depth')
44 plt.ylabel('Mean CV Score')
45 plt.title('Grid Search Scores by Max Depth')
46 plt.grid(True)
47 plt.show()

```

Best parameters: {'max_depth': 40}

Best cross-validation score: 0.852313151740708



- A max_depth of 40 was used. This was a good balance for model accuracy and computational cost.

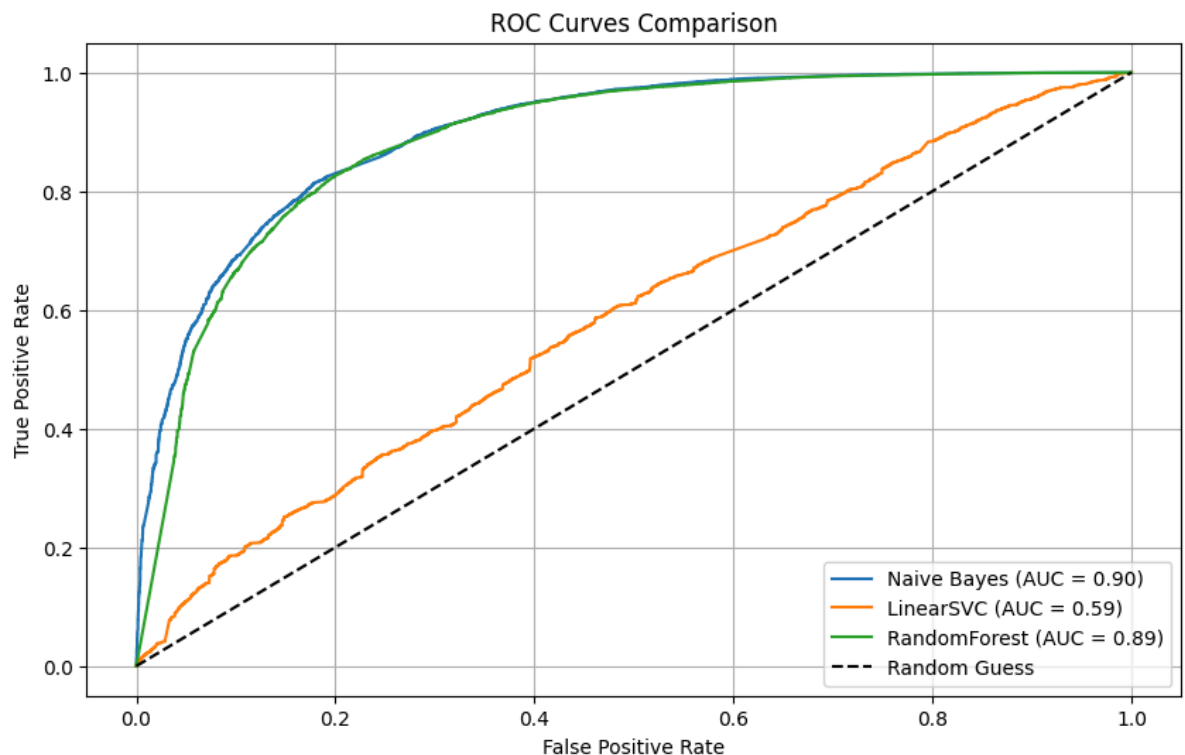
ROC Curve

- Comparing the three classifiers with ROC curve but using One vs Rest (OvR) method. This method compares one class with others by reducing the multiclass classification to multiple binary classification.

```

In [ ]: 1 from sklearn.metrics import roc_curve, auc
2
3 # Compute ROC curves and ROC AUC for Multinomial Naive Bayes
4 fpr_nb, tpr_nb, _ = roc_curve(y_test, nb_classifier.predict_proba(X_test_t
5 roc_auc_nb = auc(fpr_nb, tpr_nb)
6
7 # Compute ROC curves and ROC AUC for each classifier
8 fpr_nb, tpr_nb, _ = roc_curve(y_test, nb_classifier.predict_proba(X_test_t
9 roc_auc_nb = auc(fpr_nb, tpr_nb)
10
11 decision_scores_svm = svm_classifier.decision_function(X_test_tfidf)
12 fpr_svm, tpr_svm, _ = roc_curve(y_test, decision_scores_svm)
13 roc_auc_svm = auc(fpr_svm, tpr_svm)
14
15 rf_classifier.fit(X_train_tfidf, y_train)
16 y_pred_rf_proba = rf_classifier.predict_proba(X_test_tfidf)[: , 1]
17 fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf_proba)
18 roc_auc_rf = auc(fpr_rf, tpr_rf)
19
20 # Plot ROC curves
21 plt.figure(figsize=(10, 6))
22 plt.plot(fpr_nb, tpr_nb, label=f"Naive Bayes (AUC = {roc_auc_nb:.2f})")
23 plt.plot(fpr_svm, tpr_svm, label=f"LinearSVC (AUC = {roc_auc_svm:.2f})")
24 plt.plot(fpr_rf, tpr_rf, label=f"RandomForest (AUC = {roc_auc_rf:.2f})")
25 plt.plot([0, 1], [0, 1], 'k--', label="Random Guess")
26 plt.xlabel('False Positive Rate')
27 plt.ylabel('True Positive Rate')
28 plt.title('ROC Curves Comparison')
29 plt.legend(loc='lower right')
30 plt.grid(True)
31 plt.show()

```



- Naive Bayes: The curve for Naive Bayes has an AUC (Area Under the Curve) of 0.90, which indicates a high level of performance in distinguishing between the positive and negative classes.
- LinearSVC: The LinearSVC curve has an AUC of 0.51, suggesting that it performs only slightly better than random guessing.
- RandomForest: The RandomForest curve has an AUC of 0.89, showing good performance, though not as high as Naive Bayes.

Recommendation System

- **Build the content-based recommendation system.**
 - We create a TF-IDF matrix from the lemmatized text.
 - Compute cosine similarity between items based on their TF-IDF vectors.
 - Define a function to get recommendations for a given book title.

This approach recommends books based on the lemmatized_title_rating text

- Calculate TF-IDF vectors for "lemmatized_title_rating". Recommend books with the highest cosine similarity to a user's preferred book.
- Keyword matching: Extract keywords from "lemmatized_title_rating". Recommend books with similar keywords to a user's preferred book.

```
In [ ]: 1 nltk.download('wordnet')
2 from nltk.stem import WordNetLemmatizer
3
4 # Initialize the WordNet Lemmatizer
5 lemmatizer = WordNetLemmatizer()
6
7 # Function to Lemmatize a List of tokens
8 def lemmatize_text(tokens):
9     return [lemmatizer.lemmatize(w) for w in tokens]
10
11 # Apply Lemmatization to the 'tokenized_text' column
12 df1['lemmatized_text'] = df1['tokenized_title_rating'].apply(lemmatize_text)
```

[nltk_data] Downloading package wordnet to /root/nltk_data...

```
In [ ]: 1 # Lemmatize tokenized_title_rating column
2 df1['lemmatized_title_rating'] = df1['tokenized_title_rating'].apply(lemma
3 df1.sample(2)
```

```
Out[52]:
```

	rating	title_rating	text	user_id	helpful_vote	main_cat
54844	5	tough and informative	i knew this would be an exciting book to read ...	AE4NUL4D3AYEGBH4AY5BIKTAKYPQ	0	Buy a k
140660	1	not in the condition i requested	i am retuning this book it was used i wanted a...	AEEMSTTOB2YKEV7JKPT6CRKVBUZQ	2	E

```
In [ ]: 1 df1.info()

<class 'pandas.core.frame.DataFrame'>
Index: 156087 entries, 0 to 299999
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   rating                                156087 non-null int64
1   title_rating                          156087 non-null object
2   text                                  156087 non-null object
3   user_id                               156087 non-null object
4   helpful_vote                          156087 non-null int64
5   main_category                         156080 non-null object
6   title_book                            156087 non-null object
7   price                                 156087 non-null float64
8   tokenized_text                        156087 non-null object
9   tokenized_title_rating                156087 non-null object
10  lemmatized_text                       156087 non-null object
11  lemmatized_title_rating                156087 non-null object
dtypes: float64(1), int64(2), object(9)
memory usage: 15.5+ MB
```

```
In [ ]: 1 # Convert each List of tokens back into a string
2 df1['lemmatized_title_rating'] = df1['lemmatized_title_rating'].apply(lamb
3
4 # Create a Tfidf Vectorizer object
5 tfidf_vectorizer = TfidfVectorizer(max_features=5000)
6 tfidf_matrix = tfidf_vectorizer.fit_transform(df1['lemmatized_title_rating
7 tfidf_matrix.shape
```

```
Out[54]: (156087, 5000)
```

Cosine similarity

- Recommendation of books based on **cosine similarity** between a given book title and other books in a TF-IDF matrix. It retrieves the 5-top most similar books and returns their titles.

```
In [ ]: 1 from sklearn.metrics.pairwise import cosine_similarity
2
3 def recommend_books(book_title, tfidf_matrix, df1, n_recs=5):
4
5     # Check if the book title exists in the DataFrame
6     if book_title not in df1['title_book'].values:
7         print(f"Book title '{book_title}' not found :()")
8         return []
9
10    # Get the index of the query book
11    book_idx = df1.index[df1['title_book'] == book_title].tolist()
12    if not book_idx:
13        print(f"No index found for book title '{book_title}'.")
14        return []
15    book_idx = book_idx[0]
16
17    # Calculate cosine similarity between the query book and all books
18    cosine_similarities = cosine_similarity(tfidf_matrix[book_idx], tfidf_
19
20    # Get indices of top n most similar books, excluding the query book it
21    similar_indices = cosine_similarities.argsort()[::-n_recs-2:-1]
22    similar_indices = similar_indices[similar_indices != book_idx]
23
24    # Extract recommended book titles from DataFrame
25    recommended_books = df1.iloc[similar_indices]['title_book'].values.tol
26
27    # Return top 5 books from the list
28    return recommended_books[:n_recs]
29
```

```
In [ ]: 1 # Test 1 for the recommendation
2 book_title = 'Lord Farleigh and Miss Frost: A Regency Romance (Clairvoir C
3 recommended_books = recommend_books(book_title, tfidf_matrix=tfidf_matrix,
4 recommended_books
```

```
Out[56]: ['Addition the Fun Way Student Workbook: Requires the Addition the Fun Way Bo
ok for Kids',
'Mind over Batter: 75 Recipes for Baking as Therapy',
'Charming Bouquets: Make-a-Masterpiece Adult Grayscale Coloring Book with Co
lor Guides',
'Become Unforgettable: 7 Strategies To Scale Your Personal Brand For Maximum
Impact',
'Mr. Washington's Granite State Vacation']
```

Hybrid Recommendation System

- Multinomial Naive Bayes is better for short texts as in our case and there more suitable for the hybrid recommendation system. We will use it with the cosine similarity to have a more


```

In [ ]: 1 from sklearn.metrics.pairwise import cosine_similarity
        2
        3 def hybrid_recommend_books(book_title, tfidf_matrix, df1, nb_classifier, n
        4
        5     # Check if book title exists
        6     if book_title not in df1['title_book'].values:
        7         print(f"Book title '{book_title}' not found :(")
        8         return []
        9
        10    # Get book index
        11    book_idx = df1.index[df1['title_book'] == book_title].tolist()
        12    if not book_idx:
        13        print(f"No index found for book title '{book_title}'.")
        14        return []
        15    book_idx = book_idx[0]
        16
        17    # Content-Based Filtering with Naive Bayes
        18    content_based_score = nb_classifier.predict_proba(tfidf_matrix[book_idx])
        19
        20    # Content-Based Filtering with Cosine Similarity
        21    cosine_similarities = cosine_similarity(tfidf_matrix[book_idx], tfidf_ma
        22
        23    # Get similar book indices (excluding query book)
        24    similar_indices = cosine_similarities.argsort()[::-n_recs-2:-1]
        25    similar_indices = similar_indices[similar_indices != book_idx]
        26
        27    # Combine Scores (weighted average)
        28    final_scores = w1 * content_based_score + (1 - w1) * cosine_similarities
        29
        30    # Top Recommendations
        31    top_n_indices = final_scores.argsort()[::-n_recs:]
        32    hybrid_recommended_books = df1.iloc[top_n_indices]['title_book'].values.
        33
        34    return hybrid_recommended_books

```

```

In [ ]: 1 # Test the hybrid recommendation
        2 n_recs = 5 # Number of recommendations desired
        3 book_title = 'A Crooked Cottage by the Sea'
        4
        5 hybrid_recommendations = hybrid_recommend_books(book_title, tfidf_matrix,
        6
        7 if hybrid_recommendations:
        8     print(f"Hybrid Recommendations for '{book_title}':")
        9     for book in hybrid_recommendations:
        10         print(f"- {book}")
        11 else:
        12     print(f"No recommendations found for '{book_title}'.")
        13

```

Hybrid Recommendations for 'A Crooked Cottage by the Sea':

- The Sandman Omnibus Vol. 1
- The Keeper of Happy Endings
- The Echo of Old Books: A Novel
- The London Séance Society: A Novel
- A Crooked Cottage by the Sea

Conclusion

- We employed Multinomial Naive Bayes and cosine similarity to measure the closeness of books in the feature space. This combination allowed us to capitalize on the strengths of both methods, resulting robust recommendations.

Recommendations

Some recommendations based on the the finals result and some of the challenges encountered:

- Consider refining the model more for even better recommendations.
 - Enrich the data with more information about the books and even users' profiles
 - Implement a feedback loop where usres can also give feedback on the book recommendations they get.
-
- Jawa, Vibhu. 2021. "Accelerating TF-IDF for Natural Language Processing with Dask and RAPIDS." RAPIDS AI. <https://medium.com/rapids-ai/accelerating-tf-idf-for-natural-language-processing-with-dask-and-rapids-6f6e416429df> (<https://medium.com/rapids-ai/accelerating-tf-idf-for-natural-language-processing-with-dask-and-rapids-6f6e416429df>). Accessed April 27, 2024.
 - Scribendi Media. "How to Get Reviews for Your Book on Amazon." Scribendi Media, Dec. 2019. Image of screenshot from Amazon app with a 5-star review. Retrieved April 27, 2024, from <https://scribemedi.com/wp-content/uploads/2019/12/How-To-Get-Reviews-For-Your-Book-On-Amazon-1024x594.jpg> (<https://scribemedi.com/wp-content/uploads/2019/12/How-To-Get-Reviews-For-Your-Book-On-Amazon-1024x594.jpg>)